



01

## AWS Professional 10-Day Comprehensive Program

REF: AWS-PRO-2026

# AWS Foundations & Developer Setup

Focus: Governance, Security, and Automation First



### Today's Agenda

- AWS Cloud & Developer Environment
- AWS CLI Deep Dive & Config
- IAM Basics (Users, Groups, Policies)
- Shared Responsibility Model



### Learning Outcomes

- ✓ Master AWS Shared Responsibility
- ✓ Configure CLI securely (SSO/MFA)
- ✓ Automate IAM user provisioning
- ✓ Understand Identity Principals

Prerequisites

AWS Account

Terminal Access

Admin Permissions





## Global Infrastructure

AWS spans **33 Geographic Regions** and **105 Availability Zones (AZs)**. Each Region is a physical location with multiple AZs designed for high availability and fault tolerance.

- ✓ AZs are isolated locations within a region with independent power/cooling
- ✓ Edge Locations (400+) deliver content via CloudFront with low latency



## Core Services & Building Blocks

Comprehensive suite of IaaS, PaaS, and SaaS offerings categorized by functionality.

- Compute (EC2, Lambda)
- Storage (S3, EBS, EFS)
- Database (RDS, DynamoDB)
- Networking (VPC, Route53)



## Well-Architected Framework

Six pillars ensuring systems are secure, high-performing, resilient, and efficient: Operational Excellence, Security, Reliability, Performance Efficiency, Cost Optimization, Sustainability.

## Enterprise Governance

### AWS Organizations

Central governance and billing for multiple AWS accounts. Enables Service Control Policies (SCPs) to enforce guardrails across the organization.

### Control Tower

Automates the setup of a multi-account environment (Landing Zone) based on best practices blueprints.

### Compliance Programs

AWS supports 143 security standards and compliance certifications including PCI-DSS, HIPAA/HITECH, FedRAMP, GDPR, FIPS 140-2, and NIST 800-171.



"Security is job zero." - AWS Core Philosophy



## Free Tier Setup – Step-by-Step Guide

Module 01: Foundations

Day 1  
Setup & Config

03

1

### Account Creation & Verification

Navigate to [aws.amazon.com/free](https://aws.amazon.com/free), enter email, credit/debit card, and phone number.

- i AWS will charge a temporary \$1.00 hold to verify identity. Ensure your card allows international transactions if applicable.

2

### Secure Root User (MFA)

Login as Root. Go to **IAM Identity Center** or Security Credentials.

Menu > Security Credentials > Multi-factor authentication (MFA) > Assign MFA

Use Google Authenticator or Authy app on your mobile device to scan the QR code.

3

### Billing Preferences & Alerts

Enable PDF invoices and Free Tier usage alerts to get notified via email when limits are approached.

Receive PDF Invoice

Free Tier Alerts



### Verification Checklist

- Root account secured with MFA
- Billing currency set correctly
- \$0.00 Budget created & active
- Region checks (Use us-east-1 for billing)






### Common Pitfalls

- ✖ **Support Plan:** Selecting "Developer" or "Business" support by mistake (costs \$29+/mo). Stick to "Basic".
- ✖ **Region Blindness:** Resources created in the wrong region (e.g., Ohio vs N. Virginia) may be forgotten and billable.
- ✖ **Active Services:** Not terminating the NAT Gateway or Elastic IP after labs.



## IAM Authorization Flow



## Key Concepts

## Principals

An entity (User, Service, App) that is authenticated and allowed to make requests.

## Policies (JSON)

Documents that define permissions.

- Identity-based (Attached to User/Role)
- Resource-based (Attached to S3/SQS)

## Trust vs Permission

## Trust Policy:

*"Who can assume this role?"*

Principal: {"AWS": "arn:aws:iam::123:user/Bob"}

## Permission Policy:

*"What can they do?"*

Action: "s3:ListBucket"



### View IAM Policy Examples

```
1  {
2    "Version" :"2012-10-17",
3    "Statement": [
4      // Statement 1: Allow Read Access to Specific Bucket
5      {
6        "Sid": "AllowScopedRead",
7        "Effect": "Allow",
8        "Action": [
9          "s3>ListBucket",
10         "s3GetObject"
11       ],
12        "Resource": [
13          "arn:aws:s3:::finance-reports", // Bucket Level
14          "arn:aws:s3:::finance-reports/*" // Object Level
15        ]
16      },
17      // Statement 2: Explicit Deny for Non-SSL Traffic
18      {
19        "Sid": "EnforceTLS",
20        "Effect": "Deny",
21        "Action": "s3:*",
22        "Resource": "arn:aws:s3:::finance-reports/*",
23        "Condition": {
```

### Policy Anatomy


- 1 **Version**  
Always use `2012-10-17`.
- 2 **Sid (Statement ID)**  
Optional ID. Useful for tracking & organization.
- 3 **Effect**  
Allow or Deny. *Implicit Deny* is default.
- 4 **Condition**  
"When does this apply?" e.g., IP range, MFA, SSL.

### Pro Tips

- ✓ **Explicit Deny** always overrides Allow. Use it for guardrails.
- ✓ Avoid "Action": "\*" in production. Be specific (Least Privilege).
- ✓ Use **Policy Variables** like \${aws:username} for dynamic scaling.



## Responsibility Scope by Service Model



## Governance &amp; Ops

## Compliance Mapping

AWS manages security *of* the cloud (Physical, Network), reducing the audit scope for compliance (PCI-DSS, HIPAA).

**Example:** For RDS (PaaS), AWS handles OS patching compliance, customer handles DB user access compliance.

## OS Patch Management

**IaaS:** You control the update schedule (yum update).

**PaaS:** AWS applies minor version patches automatically during maintenance windows.

## Encryption Keys

Customer owns the keys (KMS CMK) and decides who can use them. AWS manages the hardware responsibility for your Data or



Feature / Tool	AWS Console Web UI (GUI)	AWS CLI Command Line	AWS SDK Code Libraries
⌚ Speed & Automation	<b>Slow / Manual</b> Click-ops is prone to human error. Cannot be automated easily.	<b>Fast / Scriptable</b> Great for ad-hoc scripts and repeatable tasks.	<b>Fastest / Integrated</b> Embedded in app logic. High performance.
⟳ Repeatability	<b>Low</b> Difficult to reproduce exact steps across regions/accounts.	<b>High</b> Scripts can be version controlled and re-run.	<b>Very High</b> Part of application codebase and build process.
🎓 Learning Curve	<b>Low (Beginner)</b> Visual, intuitive discovery of features.	<b>Medium</b> Requires knowing commands and parameters.	<b>High (Developer)</b> Requires programming knowledge (Python, JS, etc.).
🛡 Auth & Security	<b>SSO / MFA / Password</b> Session based. Auto-logout.	<b>Access Keys / SSO</b> Keys stored locally. Risk of leakage if not managed.	<b>IAM Roles</b> Best practice. Uses instance metadata or env vars.
💡 Primary Use Case	Initial exploration Billing & Cost Mgmt One-off diagnostics	CI/CD Pipelines Shell Automation Quick bulk operations	Application Logic Custom Tools Complex workflows

Console

CLI

SDK



# Why Use AWS CLI? Enterprise Scenarios

Module 01: Foundations

Day 1  
Foundations

08



## Automation & Scale

Critical for managing large-scale environments where Console actions are inefficient.

- ✓ Rapidly configure baselines (VPC, IAM, Config) for new accounts without manual clicks.
- ✓ Iterate through thousands of resources (e.g., tagging instances, cleaning up S3 buckets) using shell loops.



## DevOps & Integrity

Integrating AWS operations directly into development pipelines and verification workflows.

- ✓ Embed CLI commands in Jenkins, GitLab, or GitHub Actions for lightweight infrastructure deployment steps.
- ✓ Scripted verification to ensure resource configurations match the expected state (e.g., security group rules).



## Repeatability & Reliability

Ensure consistent execution across environments.

- ✓ Write scripts that handle existence checks to safely run multiple times without side effects.

## Power User Capabilities

### Demo Scenarios

For architects and leads, the CLI is faster for "live demos" or rapid prototyping than writing full SDK code for one-off tests.

### Shell Integration

Pipe output directly to tools like `jq` `grep` `awk`  
`aws s3 ls | grep "2024" | awk '{print $4}'`

### JMESPath Querying

Server-side filtering prevents fetching unnecessary data, reducing latency and cost.  
`--query 'Reservations[]Instances[]InstanceId'`



# Configure CLI: Access Keys, Region, Profiles

Module 01: Foundations

Day 1  
Developer Tools

09

1

## Standard Configuration vs. SSO

Choose your authentication method. SSO is recommended for enterprise environments.

### Method A: Static Keys

```
aws configure
```

Requires Access Key ID & Secret Key

### Method B: SSO (Recommended)

```
aws configure sso
```

Auto-rotates credentials via portal

2

## Using Named Profiles

Isolate environments (Dev, Prod) by appending `--profile` to commands.

```
$ aws configure --profile prod-user  
# Set region to us-west-2, output to json
```

*Usage:* aws s3 ls --profile prod-user

3

## Configuration Files Location

AWS CLI stores settings in two separate files in your home directory . `~/.aws/`

### config

Region, output format, sso\_start\_url

### credentials

Access Key ID, Secret Access Key  
(Sensitive!)



## Verification Checklist



Verify Identity  
`aws sts get-caller-identity`



Check Local Files  
`cat ~/.aws/credentials`



List Buckets (Test Access)  
`aws s3 ls`



## Common Pitfalls

- ✖ **Committing Secrets:** Never commit `.aws/credentials` or `.env` files to Git. Use `.gitignore`.
- ✖ **Region Mismatch:** CLI defaults to `us-east-1` if not set. Resources in Ohio (`us-east-2`) won't show up.
- ✖ **Hidden Env Vars:** Troubleshooting why profile settings aren't working? Check if `AWS_ACCESS_KEY_ID` is set in your shell.



```
> aws s3 ls

1 # 1. General Syntax Pattern
2 # aws <service> <operation> [parameters]
3 aws s3 ls
4
5 # 2. Advanced Query: Filter Running Instances & Format Output
6 aws ec2 describe-instances \
7 --filters "Name=instance-state-name,Values=running" \
8 --query "Reservations[*].Instances[*].[InstanceId, PublicIpAddress, Tags[?Key=='Name'].Value[0]]"
9 --output table
10
11 # 3. Pagination & Limits
12 aws s3api list-objects-v2 \
13 --bucket enterprise-logs-archive \
14 --max-items 50 \
15 --starting-token "eyJOZXh0VG9rZW4iOjBudWxsfQ=="
16
17 # 4. Security: Using Specific Named Profiles
18 aws iam list-users \
19 --profile prod-admin \
20 --region us-west-2
21
22 # 5. Debugging: Dry Run
23 aws ec2 run-instances --dry-run --image-id ami-12345678
```

## Command Anatomy

### 1 Service

The AWS service namespace (e.g., ec2, s3)

### 2 Operation

The API action to perform (e.g., describe-instances)

### 3 Parameters

Options starting with - or -- (filter or configure)

## Pro Tips

✓ **Server-Side Filtering:** Use --query (JMESPath) to filter data on AWS side, reducing bandwidth.

✓ **Safe Switching:** Always use --profile or export AWS\_PROFILE to prevent accidental prod commands.

✓ **Autocomplete:** Enable with complete -C '/usr/bin/aws\_completer' aws.



## Console vs CLI vs SDK: Decision Matrix

Module 01: Foundations

Day 1

Selection Strategy

11

### ⚖️ Technical Assessment Criteria

FACTOR	CONSOLE (GUI)	CLI (SHELL)	SDK (CODE)
Frequency	Ad-hoc / One-time	Recurring / Batch	Continuous / High-freq
Complexity	Simple / Visual	Moderate / Scripted	High / Logic-heavy
Audit Trail	CloudTrail (User)	CloudTrail (User/Role)	CloudTrail (App Role)
Scale	Single Resource	10s - 100s Resources	Unlimited / Parallel
Latency	High (Human speed)	Medium (Process overhead)	Low (Native integration)
Governance	Hard to enforce limits	Guardrails via wrappers	Embedded compliance logic
Dependency	Browser only	Python/Installer	Language Runtime
Error Handling	Visual Feedback	Exit Codes / Text	Try-Catch / Exceptions

### ✖️ Selection Logic



#### Scenario: Visual Verification

"I need to check why an EC2 instance isn't starting and look at screenshots."

[Use Console](#)



#### Scenario: Ops Automation

"I need to rotate keys for 50 IAM users and generate a report CSV."

[Use CLI](#)



#### Scenario: App Integration

"My web app needs to upload user avatars to S3 and generate presigned URLs."

[Use SDK](#)



*Pro Tip: Start exploring with Console, automate ops with CLI, build products with SDK.*

### Enterprise Recommendation

For production environments, restrict **Console** access. Mandate **IaC** (**Terraform/CloudFormation**) for provisioning and **SDKs** for application logic to ensure reproducibility and auditability.



▶ Development best practices

```
1 # 1. PREFERRED: Configure SSO (Avoids long-lived static keys)
2 # Interactive setup for AWS IAM Identity Center
3 aws configure sso
4
5 # Example Prompts & Inputs:
6 # SSO session name (Recommended): my-sso
7 # SSO start URL: https://my-org.awsapps.com/start
8 # SSO region: us-east-1
9
10 # 2. Switch Contexts safely using Environment Variables
11 export AWS_PROFILE="dev-profile"
12 export AWS_REGION="us-west-2"
13
14 # 3. Verify Identity (Check you are not using Root!)
15 aws sts get-caller-identity
16 # Output should show: "Arn": "arn:aws:sts::...:assumed-role/..."
17
18 # 4. ANTI-PATTERN: DO NOT DO THIS
19 # Never hardcode secrets in scripts or export them manually
20 export AWS_ACCESS_KEY_ID ="AKIAIOSFODNN7EXAMPLE" # <-- BAD!
21 export AWS_SECRET_ACCESS_KEY ="wJalrXUtnFEMI/K7MDENG/bPxRficyEXAMPLEKEY"
22
23 # 5. Handling Secrets for Apps
```

## Core Principles

### 1 SSO > Static Keys

Use Identity Center. Short-term creds auto-rotate.

### 2 Least Privilege

Don't use Admin. Scope permissions to exact needs.

### 3 MFA Everywhere

Enforce MFA for Console and CLI access.

### 4 CloudTrail

Ensure all API actions are logged for audit.

## Critical Warnings

✖ NEVER commit `env` files or credentials to Git. Use `gitignore`.

✖ Do not use the **Root User** for daily tasks. Lock it away.

ⓘ Use tools like `git-secrets` ↗



## Mini Project: IAM User Automation

Module 01: Foundations

Day 1  
Hands-on Lab

13



### IAM Provisioner

⌚ 45 Mins ⚡ Intermediate

Create a robust shell script to onboard new employees by parsing a CSV file and provisioning their IAM resources automatically.

Bash/Python

AWS CLI

IAM

CSV


### Logic Flow



users.csv



Loop & Logic



### Objectives

- Parse a `users.csv` file containing names and departments.
- Create IAM Users idempotently (check if exists first).
- Assign users to groups (Devs, Ops) based on department.
- Create initial login profile with forced password change.

>\_ provision\_users.sh

Bash

```
#!/bin/bash
# Usage: ./provision_users.sh users.csv
# CSV Format: username,group
INPUT_FILE=$1
while IFS=, read -r user group; do
    # 1. Check if user exists
    aws iam get-user --user-name "$user" &>/dev/null
    if [ $? -eq 0 ]; then
        echo "User $user already exists. Skipping."
    else
        # 2. Create User
        aws iam create-user --user-name "$user"
        # 3. Add to Group
        aws iam add-user-to-group --user-name "$user" --group-name "$group"
        # 4. Create Login Profile
    fi
done < $INPUT_FILE
```

# Terraform Automation & KMS Resource Creation

Focus: Infrastructure as Code (IaC) & Security Automation



## Today's Agenda

- IaC Concepts & Patterns
- Terraform Workflow (Init to Apply)
- Variables & Precedence Strategy
- AWS KMS Resource Creation
- Multi-Environment Delivery



## Learning Outcomes

- ✓ Implement Declarative Provisioning
- ✓ Master Terraform State Management
- ✓ Automate Security Keys (KMS)
- ✓ Manage Multi-Env Deployments





## Repeatability & Consistency

Eliminate manual configuration errors ("ClickOps") and ensure environments (Dev, Stage, Prod) are mathematically identical.

- ✓ **Idempotency:** Apply the same code multiple times without side effects.
- ✓ Disaster Recovery becomes a "terraform apply" operation.



## Reviewability via VCS (GitOps)

Infrastructure is treated as software code, residing in version control systems like Git/CodeCommit.

- ✓ Enforce peer reviews (PRs) before infrastructure changes occur.
- ✓ Full audit trail of *who* changed *what* and *when* via Git history.



## Modularity & Reusability

Encapsulate complex architecture into reusable modules to standardize patterns across the enterprise.

- ✓ Centralize best practices (e.g., "Hardened VPC Module").
- ✓ Write once, deploy everywhere (DRY principle).

## Enterprise Implications

### Drift Reduction

Detects unauthorized manual changes ("drift") by comparing the live environment state against the code. Terraform can report on or automatically revert these changes to maintain integrity.

### Policy-as-Code

Integrate tools like **Sentinel** or **OPA** (Open Policy Agent) to enforce compliance rules automatically during the plan phase (e.g., "Prohibit S3 buckets without AES-256 encryption").

### Lifecycle Management

Manage the entire lifecycle of resources, from provisioning to updates and eventual decommissioning (destroy), ensuring no "zombie" resources remain.



## Manual vs Declarative Provisioning

Module 02: Infrastructure as Code

Day 2  
IaC Concepts

16

### Manual "Click-Ops"



Admin



Configuration Drift

### Declarative (Terraform)



```
resource "aws_s3" {  
  bucket = "app"  
}
```

vs



Config

Desired State Engine



PLAN

Easy Rollback



State



Prod-01



RDS



S3

Version controlled (Git)

Dependency resolution

### Key Concepts

#### Desired State

You define **WHAT** the infrastructure should look like, not **HOW** to build it step-by-step.

Goal: 3 EC2 Instances Current: 1  
EC2 Instance Action: Add 2  
instances  
Action: Add 2 instances

#### Dependency Graph

Terraform automatically calculates the order of operations based on resource references.

e.g., VPC must exist before Subnet,  
Subnet before EC2.

#### Rollback Story

In IaC, a "rollback" is just applying a previous version of the code committed to Git.

git revert → terraform apply

#### Process




## Version-Controlled Deployments (Git Workflow)

Module 02: Infrastructure Automation

Day 2  
IaC Workflow

17

### Infrastructure CI/CD Pipeline



● Local ● CI Pipeline ● CD Pipeline

pipeline.yml

#### Module Versioning

Production infrastructure should never use or branch for modules latest main

```
module "vpc" {  
  source = "git....vpc.git?ref=v1.2.0"  
}
```

Use Semantic Versioning (Major.Minor.Patch)

#### CI Gate Checks

**Terraform Fmt**  
Ensures consistent style (HCL standards).

**Terraform Validate**  
Checks syntax and internal consistency.

**Terraform Plan**  
Speculative execution to visualize changes in PR.

**⚠ Security Note:** Ensure the CI runner has minimal IAM permissions (Least Privilege) to




## Desired State & Drift Management

Module 02: Terraform IaC

Day 2  
IaC Concepts

18

### Terraform State Reconciliation Loop



### Concepts & commands

#### ⟳ Reconciliation

Terraform compares **Desired State** (code) against **State File** and **Actual State** (Cloud) to generate a plan.

#### 🔒 State Locking

Prevents concurrent operations (e.g., two developers applying at once) using a DynamoDB table. Essential for teams.

#### ➡ Drift Management

Occurs when resources are modified outside Terraform (ClickOps).

```
# Detect drift only
terraform plan -refresh-only
```

```
# Accept drift to state
terraform apply -refresh-only
```




## Terraform Workflow Lifecycle

Module 02: Infrastructure Automation

Day 2  
Core Workflow

19

### Standard Lifecycle Operations



### Workspaces

Isolate state files for different environments without changing code structure.

```
$ terraform workspace new dev  
Created and switched to workspace  
"dev"!
```

default dev prod

### Plan Output Legend

**Create**  
+ Resource will be added.

**Update in-place**  
~ Modification of existing resource.

**Destroy**  
- Resource will be removed.



## Terraform File Structure & Modules

Module 02: Infrastructure Automation

Day 2

IaC Architecture

20

### Project Layout

```
root_module/
  main.tf
    Resources & Modules
  variables.tf
    Input Declarations
  outputs.tf
    Return Values
  versions.tf
    Constraints & Plugins
modules/
  vpc/
  security/
```

### Functional Workflow

#### providers.tf

Configures AWS region, aliases, and authentication methods.

#### versions.tf

Pins Terraform core version and provider versions (e.g., aws ~> 5.0).

### INPUTS

variables.tf  
Type Definitions

### LOGIC

main.tf  
EC2 S3 VPC

### OUTPUTS

outputs.tf  
IPs, ARNs, IDs

Remote Modules  
Registry / Git

source =  
"..."

terraform.tfstate  
Managed Automatically

### Standard Pattern

#### The Root Module

```
terraform init
terraform apply
```

#### Key Files

##### main.tf

Primary entry point. Contains resource definitions and module blocks.

##### variables.tf

Declarations for input variables. Defines types, descriptions, and defaults.

outputs.tf Values returned to the CLI or passed to other resources/modules.

#### Best Practices

- Split complex main.tf into logical files (e.g., network.tf, storage.tf)
- Always pin provider versions in versions.tf




## Terraform Variables: Types with Examples

Module 02: Infrastructure as Code

Day 2

Variable Definitions

21



```
1 variable "aws_region" {
2   type = string
3   default = "us-east-1"
4   description = "Target deployment region"
5 }
6
7 variable "instance_count" {
8   type = number
9   default = 2
10  validation {
11    condition = var.instance_count > 0 &gt; 0
12    error_message = "Instance count must be at least 1."
13  }
14 }
15
16 variable "enable_monitoring" {
17   type = bool
18   default = true
19 }
20
21 variable "resource_tags" {
22   type = map(string)
23   default = {
```



### Type System



#### Primitives

Simple values: string number  
bool



#### Collections

Grouped values: list map set



#### Structural

Complex types: object tuple



### Best Practices

- ✓ Always define **type** constraints to catch errors early during *plan*.
- ✓ Use **validation** blocks for custom rules (e.g., enforce instance sizes).
- ✓ Mark passwords/keys as **sensitive = true** to hide from logs.



## Terraform Variable Precedence

Module 02: Infrastructure as Code

Day 2  
Automation

22

### Precedence Hierarchy



### Scenarios & Patterns

#### Secrets Handling

Never commit `terraform.tfvars` to Git.

**Best Practice:** Use Env Variables (`TF_VAR_db_pass`) or fetch from AWS Secrets Manager dynamically.

#### Workspaces

When using Terraform Workspaces (dev, stage, prod), Terraform looks for specific variable definitions.

```
terraform plan -var-file="prod.tfvars"
```

#### Override Files

Files named `*_override.tf` override portions of the configuration, not just variables. Use sparingly for debugging.

Source If the same variable is defined in both source and override files, the value from the override file will be used. Commit?




## AWS KMS Resource Creation

Module 02: Infrastructure as Code

Day 2

Terraform & KMS

23



```
1 resource "aws_kms_key" "app_key" {
2   description = "KMS key for App Data encryption"
3   deletion_window_in_days = 10
4   enable_key_rotation = true
5
6   policy = jsonencode ({
7     Version = "2012-10-17"
8     Statement = [
9       Sid = "Enable IAM User Permissions"
10      Effect = "Allow"
11      Principal = {AWS = "arn:aws:iam::${var.account_id}:root" }
12      Action = "kms:/*"
13      Resource = "*"
14    ]
15  })
16 }
17
18 # Create an alias for easier reference
19 resource "aws_kms_alias" "app_key_alias" {
20   name = "alias/app-data-key"
21   target_key_id = aws_kms_key.app_key.key_id
22 }
```



### Resource Components

#### 1 enable\_key\_rotation

Essential for compliance. Rotates backing key material annually automatically.

#### 2 jsonencode()

Terraform function to render the policy. Cleaner than heredoc syntax.

#### 3 aws\_kms\_alias

Creates a display name (e.g., alias/my-key) for easier CLI/SDK usage.



### Security Pro Tips

✓ **Root Access:** Always allow the root user in the key policy to avoid creating an unmanageable key.

✓ **Deletion Window:** Default is 30 days. Set shorter (e.g., 7 days) for non-prod environments.

✓ **Outputs:** Output the arn or key\_id so dependent modules (like RDS/S3) can reference it.



## Hands-on: Provision AWS Resources with Terraform

Module 02: Infrastructure Automation

Day 2  
Lab 02

24

### Lab Details

Est. Time

30 min

Level

Intermediate

### Objectives

- Configure S3 + DynamoDB Backend for state locking
- Run full lifecycle: `init` `plan` `apply`
- Simulate & Resolve state lock conflict

### Prerequisites

Terraform v1.5+

AWS CLI v2

## 1 Configure Remote Backend (S3 & DynamoDB)

```
backend.tf
terraform { backend "s3" { bucket = "my-terraform-state-prod" }}
```

### Why Locking?

DynamoDB locking prevents concurrent operations (e.g., two developers applying at the same time) which could corrupt the state file.

#### Pro Tip:

Ensure the DynamoDB table has a partition key named `LockID`

## 2 Initialize, Plan, & Apply

```
bash - zsh
terraform init
Initializing the backend...
Terraform has been successfully initialized!
terraform plan -out=tfplan
Plan: 3 to add, 0 to change, 0 to destroy.
Saved the plan to: tfplan
terraform apply "tfplan"
aws_kms_key.example: Creating...
```



## Mini Project: Multi-Environment Infra

Module 02: Infrastructure as Code

Day 2  
Hands-on Lab

25



### Terraform Workspaces

⌚ 60 Mins ⚡ Advanced

Architect a scalable infrastructure pipeline using Terraform Workspaces to deploy distinct Dev, Staging, and Production environments from a single codebase.

Terraform

Workspaces

S3 Backend

CI/CD

### Deployment Flow



Modules + tfvars



Select Workspace

Dev (Small)

Stage (Mirror)

Prod (HA)

### Objectives

- Configure remote state in S3 with DynamoDB locking.
- Create shared modules for VPC and EC2 resources.
- Implement `.tfvars` for environment differences.
- Deploy concurrent environments using Workspaces.

### Project Structure & Logic

HCL

```
root/
  └── modules/
    └── vpc/
      ├── main.tf
      ├── outputs.tf
      ├── dev.tfvars
      └── prod.tfvars
    └── backend.tf
```

```
# main.tf: Dynamic instance count based on workspace
locals {
  env_name = terraform.workspace
  instance_counts = {
    default = 1
    prod = 3
  }
}
module "app_cluster" {
  source = "./modules/ec2"
  count = lookup(local.instance_counts, local.env_name, 1)
  tags = {
    Environment = local.env_name
  }
}
```



## Key Takeaways & Best Practices

Module 02: Infrastructure Automation

Day 2  
Summary

26



### Core Concepts & Best Practices

#### State Security



Never commit state files to Git. Use Remote Backends (S3) with KMS encryption and DynamoDB locking.

```
backend "s3" { encrypt = true }
```

#### Drift Management



Avoid "ClickOps". Manual changes in the Console cause drift. Run ~~regularly to detect differences~~ `terraform plan`.

Console changes will be overwritten

#### Version Pinning



Pin provider and module versions to prevent breaking changes in production pipelines.

```
version = "~> 3.0"
```

#### Modular Design



Don't build monoliths. Break infrastructure into reusable modules (VPC, App, DB) with clear inputs/outputs.

```
source = "./modules/vpc"
```



#### Common Pitfalls

##### Avoid These Patterns

###### Hard-Coded Secrets

Never put AWS keys or DB passwords in plain text .tf files.

**Fix: Use AWS Secrets Manager**

###### Monolithic State Files

Managing all Prod resources in one state file makes plan slow and risky (high blast radius).

**Fix: Split by Env / Layer**

###### Unpinned "Latest" Versions

Using latest means your infra might break on a random Tuesday when a provider updates.

**Fix: Use .tf.lock.hcl**

*"Infrastructure as Code is only as good as the discipline maintaining it."*

# AWS SDK & S3 Deep Dive

Focus: Infrastructure Automation & Object Storage at Scale



## Today's Agenda

- AWS SDK Automation (Boto3/Node.js)
- S3 Storage Classes & Intelligent Tiering
- Encryption, Lifecycles & Security
- Mini Project: Automated S3 Backup



## Learning Outcomes

- ✓ Script resource creation programmatically
- ✓ Architect cost-optimized storage tiers
- ✓ Implement robust data security at rest
- ✓ Automate lifecycle management policies





## Core Components: Clients, Resources & Paginator

The SDK provides low-level 1:1 API mappings (Clients) and high-level object-oriented abstractions (Resources, e.g., in Boto3).

- ✓ **Paginator:** Automatically handle `NextToken` rate over thousands of resources without manual loop logic.
- ✓ **Waiters:** Poll resource states (e.g.,) by `instance_running` until a desired state is reached.



## Resilience: Retries & Timeouts

SDKs implement built-in retry logic for throttled requests (HTTP 429) and server errors (HTTP 5xx).

- Exponential Backoff
- Connect Timeouts
- Jitter (Randomness)
- Read Timeouts



## Default Credential Provider Chain

SDKs automatically look for credentials in a specific order, enabling seamless transitions from local dev to production.

1. Environment Variables (AWS\_ACCESS\_KEY\_ID) 2. Shared Config (aws/config) 3. Container Provider (ECS)



## IAM Scoping & Least Privilege

Never hardcode credentials. Use the default credential provider chain to leverage IAM Roles (Instance Profiles/IRSA).

- ✓ **Scope Permissions:** Grant access only to specific resources (ARNs) and actions necessary for the task.
- ✓ **Condition Keys:** Enforce aws:SourceIp, aws:RequestedRegion, or aws:PrincipalOrgID.



## Secret Loading & KMS Integration

Use AWS Secrets Manager or SSM Parameter Store for configuration, not environment variables.

- ✓ **Envelope Encryption:** Use KMS for client-side encryption before data leaves the application.
- ✓ **TLS Enforcement:** Ensure SDK clients are configured to use TLS 1.2+ for all API calls.



## Resilience: Timeouts & Retries

Configure clients to fail fast and retry intelligently to prevent resource exhaustion.

- ✓ **Exponential Backoff + Jitter:** Prevent "thundering herd" problems by randomizing retry intervals.
- ✓ **Connection Timeouts:** Default timeouts are often too long. Tune connect\_timeout and read\_timeout.

## Resilience Patterns

### Exponential Jitter Formula

$$\text{Wait} = \min(\text{Cap}, \text{Base} * 2^{\text{Attempt}}) + \text{Random}$$

Adding randomness (jitter) desynchronizes client retries, flattening traffic spikes during recovery.

### Throttling Exception Handling




SDKs handle 5xx errors automatically, but ProvisionedThroughputExceededException (DynamoDB) or RequestLimitExceeded (EC2) require custom retry logic in high-throughput apps.

### Common Pitfall: Default Config

Default SDK HTTP clients often have infinite or very long timeouts. This can hang threads indefinitely if a downstream service stalls. Always set explicit timeouts.



## Storage Architecture Comparison



## Decision Guide

## When to use EBS (Block)

Think "Hard Drive". Use when you need a file system for an OS, database engine (RDS/MySQL), or low-latency disk operations.

Access: Single EC2 instance (usually)

## When to use EFS (File)

Think "Network Share". Use for content management systems, shared code repositories, or home directories across a fleet.

Access: Thousands of EC2s concurrently

## When to use S3 (Object)

Think "Internet Storage". Use for media assets, backups, data lakes, and static website hosting.

Access: Via API / HTTP (Anywhere)



## S3 Storage Classes Deep Dive

Module 03: Storage Management

Day 3

Cost & Perf Matrix

33

Metric	Standard Frequent Access	Intelligent Auto-Tiering	Standard-IA Infrequent	Glacier Flexible	Deep Archive Compliance
⌚ Use Case	Active Data	Unknown Patterns	Long-lived Data	Archive	Cold Archive
⌚ Latency	Milliseconds	Milliseconds	Milliseconds	Mins - Hours	12 - 48 Hrs
⏳ Min Duration	None	None*	30 Days	90 Days	180 Days
💰 Retrieval Fee	None	None	per GB	per GB	per GB
💸 Storage Cost	\$\$\$\$\$ Highest	\$\$\$\$ Auto-Optimized	\$\$\$\$ ~50% of Std	\$\$\$\$ ~20% of Std	\$\$\$\$ Lowest (~5%)

### ⚡ Performance First

Use **Standard** for hot data. Use **Intelligent-Tiering** if access patterns are unknown—it automatically moves objects between frequent and infrequent access tiers without performance impact or operational overhead.

### 💸 Cost Optimization

**Standard-IA** saves money for data accessed < 1/month but requires instant access. Be careful with small objects (< 128KB) as minimum object size charges apply for IA and Archive classes.

### ↳ Archival Strategy

**Glacier** and **Deep Archive** offer the lowest cost but introduce retrieval times from minutes to 48 hours. Ensure your RTO (Recovery Time Objective) aligns with the retrieval latency.



## S3 Encryption: SSE-S3 vs SSE-KMS


Module 03: Security & Storage

Day 3

Encryption Flow

34

### Envelope Encryption Process (Write)



### Key Management

#### SSE-S3 (Default)

Keys managed by S3  
Free of charge  
Transparent rotation  
Good for general data

#### SSE-KMS

Keys managed by KMS  
Supports Audit (CloudTrail)  
Separate permission control  
Customer Managed Keys (CMK)

### Performance & Limits

#### KMS API Quotas

SSE-KMS consumes KMS Request quota (e.g., 5,500 - 30,000 req/sec depending on region).

### Bucket Keys

Enable "S3 Bucket Keys" to reduce KMS calls by 99% and lower costs.

### Cost Impact:

SSE-S3 is free. SSE-KMS charges



### Bucket Policy Examples

```
1  {
2    "Version": "2012-10-17",
3    "Statement": [
4      // Statement 1: Deny Access Outside AWS Organization
5      {
6        "Sid": "DenyExternalOrgAccess",
7        "Effect": "Deny",
8        "Principal": "*",
9        "Action": "s3:*",
10       "Resource": "arn:aws:s3:::corporate-data/*",
11       "Condition": {
12         "StringNotEquals": {
13           "aws:PrincipalOrgID": "o-a1b2c3d4e5"
14         }
15       }
16     },
17     // Statement 2: Allow Public Read from Specific Referer
18     {
19       "Sid": "AllowWebAssetsHotlinking",
20       "Effect": "Allow",
21       "Principal": "*",
22       "Action": "s3:GetObject",
23       "Resource": "arn:aws:s3:::public-assets/*"
24     }
25   }
26 }
```



### Condition Keys

#### 1 aws:PrincipalOrgID

Restricts access to principals belonging to a specific AWS Organization ID.

#### 2 aws:Referer

Limits access based on the HTTP Referer header. Useful for web assets.

#### 3 Deny + StringNotEquals

Creates a strict data perimeter. "If you are NOT in my Org, DENY."



### Security Pro Tips

⚠ **Referer Spoofing:** aws:Referer can be spoofed. Never use it for private/sensitive data.



**Block Public Access (BPA):** Account-level BPA overrides these policies. Enable BPA for all private buckets.



✓ Use Service Control Policies (SCPs) to enforce aws:PrincipalOrgID across accounts.



## Versioning & Lifecycle: Transitions/Expiration


Module 03: AWS SDK + S3 Deep Dive

Day 3

S3 Management

36

### Object Lifecycle Policy Flow



### Version Stack Behavior

What happens when you overwrite?



i Lifecycle rules can target **current versions**, **non-current versions**, or **delete markers**.

### Why Versioning?

Versioning enables you to recover objects from accidental deletion or overwrite.

- ✓ Protect against accidental deletes
- ✓ Required for Cross-Region Replication (CRR)

### Delete Markers

"A delete marker is a placeholder for a deleted object. It has no data content."

**Behavior:** If you GET an object with a delete marker as the current version, S3 returns 404.

**Cleanup:** Use lifecycle rules to remove "expired object delete markers" to save costs.

### Replication Requirement

Both source and destination buckets **must** have versioning enabled for replication to work.



## S3 Storage Class Selection Logic

Data Access Pattern?

Intelligent-Tiering

S3 Standard

Standard-IA

Deep Archive

~12-48 hrs, lowest cost



Unknown / Changing?

Predictable?

Frequently Accessed?

Min. Duration Costs

Standard

None

Intelligent-Tiering

None\*

Standard-IA

30 Days

Glacier Flexible

90 Days

Deep Archive

180 Days

\*Monitoring fee applies per 1k objects

! Retrieval Cost Warning: accessing data from IA or Glacier tiers incurs a per-GB retrieval fee.



## Hands-on: CLI Scripting & SDK S3 Configuration

Module 03: AWS SDK + S3 Deep Dive

Day 3  
Lab 03

38

### Lab Details

Est. Time 45 min

Level Advanced

### Objectives

- Provision EC2 via Python SDK (Boto3)
- Create S3 Bucket & Enable Versioning
- Apply Lifecycle Policy (Transition to Glacier)
- Validate via CLI & Clean Up

### Prerequisites

Python 3.9+

Boto3 Lib

AWS CLI v2

### 1 Provision EC2 Instance via Boto3

```
launch_ec2.py\n\nimport boto3      # Initialize EC2 Client\nec2 = boto3.     client ('ec2' , region_name= 'us-east-1') # Launch t3.micro instance with tags
```

### 2 Configure S3 (Versioning & Lifecycle)

```
configure_s3.py\n\nimport boto3 s3 = boto3.     client ('s3' ) bucket = 'corp-archival-lab-day3-UserXYZ'
```

#### Parameter Note

#1. *Enable Versioning*: Ensure your bucket name is globally unique. Append your username or a random ID.

#### Cost Optimization

Transitioning to Glacier significantly reduces costs for rarely accessed data (\$0.0036/GB) vs Standard (\$0.023/GB).





## Case Study: Media Archival System

Module 03: AWS SDK & S3 Deep Dive

Day 3

Intelligent-Tiering

39



### Business Outcomes

#### Total Savings

**40%** reduction in TCO

#### Operational Overhead

**Zero** manual lifecycle changes

#### Durability

**99.99... %** (11 9's) across all tiers

#### lifecycle-config.xml

```
<LifecycleConfiguration> <Rule> <ID>MoveToIntelligentTiering</ID> <Filter> <Prefix>raw/</Prefix>
```



## Mini Project: Automated S3 Backup

Module 03: Storage & Archival

Day 3  
Hands-on Lab

40

### S3 Lifecycle Automator

⌚ 60 Mins ⚡ Intermediate

Design a cost-optimized backup solution that automatically transitions aging data from standard storage to cold archival storage using Terraform.

Terraform S3 Glacier KMS

### Data Lifecycle Flow



Day 0: Upload  
Standard Class

30 Days ➔



Standard-IA  
Cost Saving

90 Days ➔



🔒 Encrypted with SSE-KMS

Glacier  
Deep Archive

### Objectives

- ✓ Provision a private S3 bucket with versioning enabled.
- ✓ Enforce Server-Side Encryption (SSE-KMS) for data security.
- ✓ Configure Lifecycle Rules: Transition to IA after 30 days.
- ✓ Archive to Glacier after 90 days and expire after 1 year.

main.tf

HCL

```
resource "aws_s3_bucket" "backup" {
  bucket = "corp-data-backup-2026"
}

resource "aws_s3_bucket_lifecycle_configuration" "archive" {
  bucket = aws_s3_bucket.backup.id
  rule {
    id = "log-archival-policy"
    status = "Enabled"
    transition {
      days = 30
      storage_class = "STANDARD_IA"
    }
    transition {
      days = 90
      storage_class = "GLACIER"
    }
  }
}
```